

SoulScope Narrative Registry v0.1

Status: Canonical scientific implementation companion — architecture final, copy validation pending

Governing authority: SoulScope Canon v1.3

Upstream: Evidence Marker Registry v0.1; Constellation Dimension Registry v0.1; State Registry v0.1; Interaction Registry v0.1; Pattern Registry v0.1

Downstream: User Result Interface

Date: August 14, 2026

NORTH STAR

The narrative explains completed reasoning. It never performs new reasoning.

The user should be able to say:

I understand what SoulScope noticed, how strongly it noticed it, what it compared, what remains uncertain, and why the result was described this way.

Outline

The Narrative Engine SHALL:

1. consume only the completed immutable Semantic Result;
2. preserve the scientific meaning of every upstream object;
3. communicate scope, evidence, uncertainty, and alternatives;
4. distinguish observation from interpretation;
5. distinguish confidence from intensity;
6. distinguish personal reference from population reference;
7. preserve unresolved findings;
8. avoid identity, diagnosis, causation, hidden-state, and certainty claims;
9. expose evidence progressively rather than overwhelming the user;
10. remain deterministic in meaning even when wording varies.

Canonical sequence:

Completed Semantic Result → Narrative Plan → Bounded Language → User Result.

1. Purpose

The Narrative Registry governs how SoulScope transforms machine-readable semantic conclusions into user-facing reflection.

It defines:

- what may be said;
- what must be said;
- what may not be said;
- how confidence affects language;
- how uncertainty is communicated;
- how Evidence is summarized;
- how comparisons are framed;
- how first scans differ from longitudinal scans;
- how unresolved results are handled;
- how patterns, states, dimensions, and interactions are translated.

The narrative layer is a translation system.

It is not a scientific scoring system.

2. Narrative authority boundary

The Narrative Engine SHALL NOT:

- inspect raw audio;
- independently interpret transcripts;
- calculate acoustic features;
- recompute Evidence;
- rescore Dimensions;
- choose States;
- choose Interactions;
- choose Patterns;
- fill missing values;
- infer D3 constructs;
- infer emotion labels;
- add causal explanations;
- identify personality traits;
- override abstention;
- make an unresolved result sound resolved.

HARD RULE

$$\text{Narrative authority} \leq \text{Semantic Result authority}$$

The narrative can simplify wording.

It cannot strengthen meaning.

3. Narrative input

The Narrative Engine receives only a completed authorized object:

```
NarrativeInput

scan_id
protocol_version

pattern_result

constellations:
  COG
  REG
  CAP
  EXP

dimensions[]

states[]
qualifiers[]

interactions[]

evidence_summaries[]

contradictions[]

missingness[]

confounds[]

reference_information

confidence_information

temporal_information
```

```
longitudinal_information  
  
prohibited_claims[]  
  
user_language_keys[]  
  
registry_versions[]
```

No raw acoustic waveform is supplied to the Narrative Engine.

4. Narrative output

The Narrative Engine produces:

```
NarrativeResult  
  
headline  
  
whole_scan_summary  
  
constellation_sections[]  
  
interaction_sections[]  
  
evidence_explanations[]  
  
uncertainty_statements[]  
  
comparison_statements[]  
  
longitudinal_statements[]  
  
alternative_explanations[]  
  
reflection_prompts[]  
  
progressive_disclosure_copy[]  
  
safety_notices[]  
  
narrative_manifest
```

5. Narrative levels

SoulScope recognizes five levels of user-facing semantic language.

N0 — Measurement

Direct description of measured signal.

Example:

“Your pitch range was broader in the second response than in the first.”

N1 — Evidence

Description of qualified acoustic structure.

Example:

“Your vocal modulation expanded when the topic changed.”

N2 — Functional Dimension

Bounded interpretation of a registered Dimension.

Example:

“That contributed to a higher Activation estimate in this scan.”

N3 — Constellation / Interaction

Description of a completed higher-order relationship.

Example:

“Your vocal pattern became more mobilized while remaining comparatively stable.”

N4 — Reflective Hypothesis

A possible lived interpretation, explicitly optional.

Example:

“One possibility is that this topic required more active engagement while you still maintained a steady way of responding.”

HARD RULE

N4 may never be presented as fact.

6. Claim-strength hierarchy

Narrative strength must decrease as semantic distance from measurement increases.

Conceptually:

$$Certainty_{N0} \geq Certainty_{N1} \geq Certainty_{N2} \geq Certainty_{N3} \geq Certainty_{N4}.$$

A high-confidence Dimension does not justify high-certainty lived-meaning language.

7. Canonical sentence architecture

Preferred user-facing explanatory structure:

Scope + Observation + Comparison + Bounded meaning + uncertainty where relevant

Example:

“In this scan, your speech flow remained comparatively continuous across all three responses. That supported a higher Focus Continuity estimate. This describes the organization of these responses, not your general ability to concentrate.”

8. Scope-first language

Narrative SHOULD begin claims with an explicit scope.

Preferred:

In this scan... Across these responses... During the second prompt...

When the topic changed...
Compared with your first response...
Compared with your recent scans...
Across the approximately 30-second response...

Avoid:

You are...
Your voice always...
You tend to...
This is how you...

unless a separately validated longitudinal statement explicitly permits recurrence language.

9. Observation before interpretation

Preferred sequence:

"Your pause load increased and your speech runs became shorter during the troubling-topic prompt. Together, those measurements contributed to greater temporal fragmentation."

Then:

"That reduced the Focus Continuity estimate for that part of the scan."

Not:

"You lost focus when talking about something troubling."

10. Evidence before meaning

Whenever practical:

measurement → Evidence → interpretation.

The narrative SHOULD be able to answer:

Why are you telling me this?

without introducing a new claim.

11. Confidence language

Confidence and Dimension direction are separate.

A high estimated Dimension with low confidence should sound different from a high Dimension with high confidence.

High confidence

Preferred:

showed
was clearly supported
consistently appeared
the evidence strongly supported

Moderate confidence

Preferred:

suggested
leaned toward
appeared to show
was moderately supported

Limited confidence

Preferred:

may have reflected
showed some evidence of
leaned slightly toward
was only weakly supported

Unresolved

Preferred:

did not resolve clearly
there was not enough consistent evidence

the available evidence was mixed
SoulScope could not reliably estimate this

12. Confidence SHALL NOT become emotional certainty

Prohibited:

"SoulScope is 92% sure you were stressed."

Even if:

confidence = 0.92

because the confidence value relates to the registered construct, not an unmeasured emotional state.

13. Intensity language

Dimension magnitude and confidence SHALL remain distinct.

Example:

posterior_mean = .82
confidence = .46

Preferred:

"The available evidence leaned toward a higher Activation pattern, but confidence was limited."

Not:

"Your Activation was definitely very high."

14. User-facing numeric policy

Raw posterior values SHOULD NOT normally be shown as psychological-style scores.

Preferred:

```
lower  
slightly lower  
near reference  
slightly higher  
higher
```

with uncertainty disclosure.

Numeric data may be available in technical/progressive disclosure views where appropriate.

15. Balanced-zone language

The balanced zone SHALL NOT be called:

```
ideal  
optimal  
normal  
healthy  
perfect  
regulated
```

Preferred:

“This point sat near the current reference range.”

or:

“The evidence was mixed around the center of the active reference.”

16. Unknown language

Unresolved results SHALL be described explicitly.

Preferred:

"SoulScope did not have enough reliable evidence to resolve this point."

"The measurements pointed in different directions, so this result remained unresolved."

"This scan does not contain the type of observation needed to estimate Recovery."

Never:

"Recovery was neutral."

"Reserve was average."

"Relational Availability was balanced."

17. Structural unresolved language

For current D3 parameters:

Recovery

"This scan does not include a dedicated recovery period, so Recovery is not estimated."

Reserve

"The current scan does not include the repeated load structure needed to estimate Reserve."

Relational Availability

"These prompts are spoken individually rather than interactively, so Relational Availability is not estimated."

These statements SHOULD appear only when the user encounters the relevant Dimension.

18. First-scan language

On the first valid scan:

personal_baseline_status = ABSENT

Narrative SHALL NOT imply a personal normal.

Preferred:

"This is your first SoulScope scan, so comparisons are based primarily on differences within today's three responses and any validated compatible reference information."

Not:

"This is higher than usual for you."

19. Within-session language

P1 is an open reference condition.

Preferred:

"Compared with your open first response..."

"When the topic moved from your open response to something troubling..."

Not:

"Compared with your neutral baseline..."

20. Population-reference language

If population reference is used:

Preferred:

"Relative to the compatible reference sample used by this version of SoulScope..."

Not:

"Compared with normal people..."

Population reference SHALL NOT imply health, desirability, or ideality.

21. Personal-baseline language

When an established personal baseline exists:

Preferred:

“Compared with your established SoulScope reference...”

“This was higher than your recent personal range...”

“This pattern sat close to what has appeared across your recent compatible scans...”

The narrative SHALL preserve baseline trust.

22. Provisional personal-baseline language

If the personal baseline is still provisional:

Preferred:

“Compared with the early personal reference SoulScope is beginning to build for you...”

The wording should communicate that the reference may still change.

23. Longitudinal language

Longitudinal narrative may describe:

- recurrence;
- movement;
- widening/narrowing range;
- changing interaction regimes;
- increased/decreased deviation from personal reference.

It SHALL NOT automatically describe:

- improvement;
- deterioration;
- healing;
- regression;

- progress;
 - decline.
-

24. Recurrence language

Preferred:

"A similar pattern has appeared in three recent compatible scans."

Not:

"This is part of your personality."

Repeated observation remains repeated observation.

25. Change language

Preferred:

"This scan showed more vocal modulation than your recent personal reference."

Not:

"You have become more expressive."

Preferred:

"The relationship between Activation and Effort Cost changed relative to your recent scans."

Not:

"You are coping better now."

26. Constellation narrative hierarchy

Each Constellation section SHOULD contain:

1. one bounded summary;
2. strongest resolved Dimensions;
3. one important contrast or Evidence explanation;

4. uncertainty or missingness if material;
 5. optional evidence-trace link.
-

27. Cognitive Form narrative

Canonical framing:

How was speech-production organization formed and sustained during these responses?

Preferred language:

organized
continuous
fragmented
reorganized
more/less interrupted
broader/narrower modulation
greater observable speech-production demand

Avoid:

clear-minded
scattered thoughts
smart
confused
focused person
ADHD
overthinking

28. Organization

Preferred:

“Your speech timing and response structure stayed comparatively organized.”

“This response showed more interruption and reorganization than the open-reference response.”

Prohibited:

"Your thoughts were organized."

"Your mind became disorganized."

29. Exploration

Preferred after validation:

"Your voice used a broader variety of modulation across the changing topics."

Not:

"You were more curious."

"You explored your feelings."

30. Focus Continuity

Preferred:

"Your speech flow remained comparatively continuous."

"There were more interruptions and shorter speech runs in this part of the scan."

Not:

"You had trouble focusing."

31. Processing Demand

Preferred after validation:

"Several timing and voice-quality measurements suggested greater speech-production demand in this response."

Not:

"Your brain was working harder."

"You were cognitively overloaded."

32. Regulatory Motion narrative

Canonical framing:

How did the vocal pattern mobilize, remain stable, and change as the speaking context changed?

Preferred vocabulary:

mobilized
stable
variable
reconfigured
context-sensitive
changed across prompts

Avoid:

emotionally stable
regulated person
dysregulated
resilient
anxious
stressed

33. Activation

Preferred:

"The vocal pattern showed greater Activation relative to your open response."

Better evidence-rich version:

"Pitch, energy, and timing shifted together in a way that supported greater acoustic Activation."

Not:

"You were more stressed."

"You were more aroused."

34. Stability

Preferred:

"Within this response, the vocal pattern remained comparatively stable."

"The pattern showed greater short-term variation."

Not:

"You were emotionally stable."

35. Flexibility

Preferred after validation:

"Your vocal pattern changed coherently as the topic changed."

Not:

"You adapt well."

"You are psychologically flexible."

36. Available Capacity narrative

Canonical framing:

What observable output was maintained during the scan, and what measurable cost accompanied it?

This Constellation SHALL remain especially scan-bound.

37. Mobilization

Preferred:

"Your vocal output became more mobilized as the topic changed."

Not:

"You became more motivated."

38. Effort Cost

Preferred after validation:

"Maintaining your response was accompanied by greater observable acoustic and timing cost."

Not:

"You were exhausted."

"You were pushing yourself."

39. Sustainability

Preferred:

"The organization of your vocal output remained comparatively maintained across the approximately 30-second response."

Not:

"You have strong endurance."

40. Expressive Interface narrative

Canonical framing:

How much acoustic variation was present, and how was that variation deployed across these contexts?

41. Range

Preferred:

"Your voice used a broader range of pitch, energy, and other modulation during this part of the scan."

Not:

"You were highly expressive emotionally."

42. Openness

Preferred after validation:

"A larger share of the available vocal range was deployed across the different contexts."

Not:

"You were emotionally open."

43. Restraint

Preferred after validation:

"A broader range was available elsewhere in the scan, while this response used comparatively less of it."

Not:

"You were holding back."

"You were repressing something."

44. State narrative

State labels remain secondary to geometry.

Preferred structure:

"Your Cognitive Form result fell within the Continuous Organization region, supported mainly by Organization and Focus Continuity."

If qualifier exists:

"Processing Demand was also higher, which adds context without changing the underlying state."

45. Boundary-blend language

Preferred:

"This result sat between two neighboring regions rather than resolving clearly into one."

Then identify both.

Example:

"The Cognitive Form geometry sat between more continuous organization and greater interruption."

Avoid:

"You are partly State A and partly State B."

46. Unresolved State language

Preferred:

"The individual measurements were informative, but they did not form one sufficiently clear State."

This communicates that abstention does not mean the scan failed.

47. Interaction narrative law

Interactions SHALL be phrased as relationships among independently observed structures.

Preferred:

“These two patterns appeared together.”

“The relationship between them changed across prompts.”

“One pattern remained maintained while another showed greater cost.”

Avoid causal language unless explicitly validated.

48. Reinforces

Preferred:

“These two parts of the result pointed in the same functional direction.”

Not:

“One reinforced the other biologically.”

49. Buffers

Preferred user translation:

“Higher stability appeared alongside less disruption than the demand pattern alone might suggest.”

Avoid:

“Your stability protected you.”

unless the registered user-language policy explicitly permits a softened metaphor after comprehension testing.

50. Amplifies

Preferred:

“The two patterns appeared together in a way that made the observable cost more pronounced.”

Not:

"One caused the other to intensify."

51. Masks

Preferred user language SHOULD generally avoid the word `Masks`.

Recommended translation:

"Your responses remained organized even while another part of the analysis showed greater observable cost."

Canonical internal relationship may remain:

`MASKS`

but the frontend user-language key SHOULD default to:

`MAINTAINED_OUTPUT_WITH_COST`

because "masking" strongly invites hidden-state interpretation.

52. Compensates

Preferred:

"A stronger pattern in one area appeared alongside maintained function despite a limitation elsewhere."

Not:

"You compensated for weakness."

53. Constrains

Preferred only with required temporal/contextual support:

"As one part of the pattern narrowed, another showed less range of movement as well."

Avoid psychological or causal wording.

54. Protects

Preferred user language SHOULD usually be softened:

“Stable organization appeared alongside lower observable cost under the more demanding context.”

Avoid:

“Your Stability protected you.”

unless future validation and comprehension testing support that wording.

55. Redirects

Preferred:

“As the context changed, one part of the vocal pattern increased while another became narrower.”

Not:

“Your emotional energy moved inward.”

56. Destabilizes

Preferred:

“The change in one part of the result appeared alongside greater variability in another.”

Only use directional wording if temporal evidence supports it.

57. Integrates

Preferred user-language key:

MULTI_DOMAIN_ALIGNMENT

Suggested wording:

"Several independently measured parts of the scan aligned with relatively little contradiction."

Avoid:

"Your system was integrated."

58. Reveals

Preferred:

"The second result added context to an otherwise ambiguous first result."

Not:

"This revealed what was really going on."

59. Shifts

Preferred:

"The relationship among the measured patterns changed as the speaking context changed."

Not:

"Your emotional state shifted from negative to positive."

60. Pattern headline law

A Pattern headline should summarize the scan, not label the person.

Preferred:

Maintained with greater observable cost

Not:

The Maintainer

Preferred:

Contextual reconfiguration

Not:

The Adapter

61. Pattern headline syntax

Preferred pattern headline construction:

[Functional structure]

or:

[Functional structure] + [one modifier]

Examples:

Coherent maintenance
Maintained with greater observable cost
Contextual reconfiguration
Variable reconfiguration
Cross-domain mismatch
Temporal relation shift

62. Pattern summary paragraph

Recommended template:

“Across this scan, **{pattern_bounded_meaning}**. This was supported most strongly by **{top_constellation_1}** and **{top_constellation_2}**, with **{interaction_or_modifier}** adding important context. **{uncertainty_statement_if_needed}**”

63. No-Pattern headline

When no Pattern is published:

Preferred:

No single whole-scan pattern resolved clearly

Supporting text:

“Several individual observations were informative, but they did not form one sufficiently supported overall structure.”

This SHALL be treated as a legitimate result.

64. Evidence explanation

Every important result SHOULD have a short “Why?” explanation.

Example:

Why this appeared: Your pause load increased, speech runs shortened, and timing became more variable in the second response. Those measurements contributed to the Focus Continuity result.

This is more important than displaying an opaque score.

65. Progressive Evidence disclosure

Recommended hierarchy:

Level 1

Plain-language Evidence:

“Your pauses became longer.”

Level 2

Technical Evidence:

Median pause duration
P1: 410 ms
P2: 620 ms

Level 3

Scientific metadata:

feature_id
extractor
version
quality
reference
rule

The main interface does not need to show all three simultaneously.

66. Explanation graph language

Trace views should permit:

Pattern
↓
Interaction
↓
Constellation
↓
Dimension
↓
Evidence
↓
Measurement

Preferred UI copy:

"See why SoulScope reached this result."

Not:

"See what your voice revealed."

67. Contradiction language

Contradiction should be understandable.

Preferred:

“Not all measurements pointed in the same direction.”

“Timing Evidence supported this result, while voice-quality Evidence was less consistent.”

“Because the Evidence was mixed, confidence was reduced.”

This is scientifically valuable and should not be hidden.

68. Alternative interpretation language

Where the Decision Ledger contains meaningful alternatives:

Preferred:

“Another plausible reading was **{alternative}**, but it was less strongly supported because **{reason}**.”

This feature may be limited to progressive disclosure.

69. Confound language

Confounds should be communicated when they materially affect the result.

Preferred:

“Background noise reduced confidence in some voice-quality measurements.”

“A microphone change limited comparison with earlier scans.”

Not:

“Your result may be wrong.”

unless the result is actually invalid.

70. Invalid scan language

Invalid acquisition is different from uncertainty.

Preferred:

“SoulScope could not produce a reliable result from this recording because the audio did not meet the measurement requirements.”

Then provide retry guidance if permitted.

71. Emotional-language boundary

The Narrative Engine SHALL NOT convert:

Activation

into:

stress
fear
excitement
anger
anxiety

unless a separately validated emotion modality is adopted in a future Canon version.

Likewise:

narrow Range

does not become:

sadness
withdrawal
suppression

72. Transcript-content boundary

Current SoulScope voice analysis does not use semantic transcript content to determine Constellation meaning.

Therefore narrative SHALL NOT say:

“Because you talked about your relationship...”

unless the system is simply identifying which prompt was administered.

The allowed form is:

“During the troubling-topic prompt...”

not:

“When discussing your relationship...”

unless the user-facing transcript layer separately permits that content.

73. Prompt emotional boundary

P2 SHALL NOT be narrated as:

negative condition
stress condition
sad condition

P3 SHALL NOT be narrated as:

positive condition
hopeful emotional condition
recovery condition

They are context prompts.

74. Personal agency language

Tier C reflection SHOULD preserve choice.

Preferred:

One possibility is...
You may recognize this as...
This could reflect...
Another explanation could be...
Only you can decide whether this fits your experience.

Avoid:

This means...
The reason is...
Your voice is telling us...
You are actually...

75. Reflective hypotheses

Reflective hypotheses are optional.

A result does not require one.

The engine may produce:

```
reflection_status = NONE
```

when the evidence does not justify a useful Tier C interpretation.

Scientific value does not depend on always generating a psychologically resonant reflection.

76. Reflection-generation rule

A Tier C reflection may be offered only when:

1. upstream result is production-active;

2. confidence meets active policy;
 3. no prohibited inference is required;
 4. at least one non-pathologizing alternative exists;
 5. the wording remains conditional;
 6. the reflection adds user value beyond restating the measurement.
-

77. Reflection alternatives

Whenever a reflective hypothesis is offered, the narrative **SHOULD** preserve alternative explanations.

Example:

“One possibility is that this topic required more active engagement. Differences in speaking style, familiarity with the topic, or how you chose to respond could also contribute.”

This does not require dumping every possible confound.

It does require avoiding a single inevitable explanation.

78. Narrative confidence tiers

Canonical narrative tiers:

NT1_CLEAR
NT2_SUPPORTED
NT3_TENTATIVE
NT4_UNRESOLVED

NT1_CLEAR

Strong production-authorized result with high evidence quality and limited contradiction.

NT2_SUPPORTED

Usable result with moderate confidence.

NT3_TENTATIVE

Limited-confidence result permitted only with explicit uncertainty wording.

NT4_UNRESOLVED

No semantic claim beyond explaining why resolution failed.

79. Narrative-template object

Every user-language template SHALL contain:

```
template_id
template_version

semantic_object_type

required_status
minimum_narrative_tier

allowed_variables[]

required_scope_phrase

required_uncertainty_phrase

optional_evidence_phrase

prohibited_terms[]

localization_key

validation_status
```

80. Lexicon governance

The Narrative Registry SHALL maintain explicit vocabularies.

Preferred functional words

```
appeared
showed
suggested
supported
```

contributed
shifted
changed
remained
expanded
narrowed
increased
decreased
organized
maintained
variable
continuous
context
comparison
evidence
uncertainty

High-risk words requiring policy review

stress
anxiety
depression
trauma
avoidance
suppression
repression
coping
resilience
dysregulation
burnout
confidence
honesty
deception
masking
healing
recovery
personality

Prohibited without future scientific authorization

diagnosis
disease
lie
truthful

psychopathology
mental disorder
soul measurement
spiritual state
hidden truth

81. Identity-language prohibition

The system SHALL NOT create constructions such as:

You are a...
You are the kind of person who...
This is your type...
Your personality is...

A Pattern, State, or Dimension describes the scan.

82. Fixed-trait prohibition

Avoid:

"You are restrained."

Use:

"This response showed comparatively less deployment of the vocal range available elsewhere in the scan."

83. Moral-language prohibition

No result may imply:

good
bad
strong person
weak person
honest
dishonest

brave
cowardly
authentic
inauthentic

84. Clinical-language prohibition

No narrative shall use clinical terminology merely because it sounds explanatory.

Examples prohibited without a validated clinical product:

dysregulated
depressed
manic
traumatized
burned out
anxious disorder
cognitively impaired

85. Physiological causation boundary

Acoustic measurements may be influenced by physiology.

SoulScope SHALL NOT convert those measurements into unmeasured physiological claims.

Avoid:

“Your nervous system activated.”

Use:

“The vocal pattern showed greater acoustic Activation.”

86. Narrative personalization

Personalization may include:

- user's name where appropriate;

- scan date;
- personal-baseline comparison;
- history;
- user-selected reflection notes.

Personalization SHALL NOT include invented personality inferences.

87. User reflection input

If the user supplies their own interpretation:

"I think I was anxious there."

SoulScope may preserve it as:

USER_REFLECTION

but SHALL NOT retroactively convert the scientific result into:

ANXIETY_DETECTED

User meaning and instrument evidence remain separate.

88. Narrative + user meaning

The product may explicitly present:

What SoulScope noticed

and separately:

What this means to you

This reinforces the foundational principle:

The instrument provides evidence. The individual provides meaning.

89. Narrative hierarchy for main result

Recommended main-result order:

1. Whole-scan headline

Pattern or no Pattern.

2. 2–4 sentence summary

Most important structures.

3. Four Constellation summaries

Brief and bounded.

4. Important Interactions

Maximum two in primary display.

5. What changed across prompts

Where supported.

6. Evidence access

“Why this result?”

7. Uncertainty / limitations

Where material.

8. Reflection

Optional.

90. Main-result length discipline

Narrative richness should come from structure, not verbosity.

Recommended initial main-result targets:

headline: 3-8 words

whole-scan summary:
50-120 words

each Constellation:
35-80 words

primary Interaction:
25-60 words

reflection:
20-60 words

Technical details belong in progressive disclosure.

91. Whole-scan example

Backend

Pattern:
MAINTAINED_WITH_COST

COG:
CONTINUOUS_ORGANIZATION

REG:
STABLE_PATTERN
+ HIGHER_ACTIVATION

CAP:
MAINTAINED_OUTPUT
+ HIGHER_EFFORT_COST

EXP:
BROAD_MODULATION

Interaction:
CROSS_DOMAIN_MAINTENANCE_COST_MISMATCH

User narrative

Maintained with greater observable cost

Across this scan, your responses remained comparatively organized and stable, even while separate timing and voice-quality evidence suggested that maintaining that output carried greater cost. Activation also rose in parts of the scan, while your overall vocal range remained broad. This describes the pattern present in these three responses rather than a fixed trait or explanation for how you felt.

92. Context-shift example

Backend

Pattern:
CONTEXTUAL_RECONFIGURATION

P1 -> P2:
Activation increases
Range narrows

P2 -> P3:
Activation decreases
Range broadens

Stability:
maintained

Flexibility:
higher

User narrative

Contextual reconfiguration

Your vocal pattern changed noticeably as the subject matter changed. During the troubling-topic response, acoustic Activation increased while expressive range became narrower. During the future-oriented response, the pattern shifted again. Because these changes occurred while short-term Stability remained comparatively strong, the scan showed coordinated context-sensitive reconfiguration rather than simply greater variability.

No emotional-state labels are required.

93. Mixed-Evidence example

Backend

```
COG-P4 Processing Demand:  
  mean high  
  confidence limited  
  
support:  
  timing  
  
contradiction:  
  spectral  
  
phonation:  
  missing
```

User narrative

"Some timing measurements suggested greater speech-production demand, but other Evidence was less consistent and one required Evidence family was unavailable. SoulScope therefore treated this result as uncertain rather than drawing a firm conclusion."

94. No-pattern example

Backend

```
Pattern:  
  NO_PATTERN_PUBLISHED
```

User narrative

No single whole-scan pattern resolved clearly

Several parts of this scan were informative, but they did not combine into one sufficiently supported overall Pattern. SoulScope has preserved the individual Constellation results rather than forcing them into a single label.

95. First-scan example

“Because this is your first compatible SoulScope scan, the result relies mainly on differences within today's three responses. As more compatible scans are collected, SoulScope may be able to compare future measurements with a more established personal reference.”

96. Longitudinal example

“Compared with your established personal reference, today's vocal range was broader and Effort Cost was lower. The relationship between Activation and Sustainability also differed from your recent scans. These are measured changes in the SoulScope pattern; they do not by themselves indicate improvement or decline.”

97. Narrative consistency

Given the same immutable Semantic Result and the same Narrative Registry version, semantic meaning SHALL remain reproducible.

Wording may vary within approved templates.

Meaning may not.

98. LLM role

A language model MAY perform:

- grammatical realization;
- approved paraphrase;
- tone adaptation;
- localization support;
- compression;
- expansion within registered limits.

It SHALL NOT perform:

- inference;
- classification;
- causal reasoning;
- diagnostic reasoning;
- Pattern selection;

- State selection;
- confidence adjustment;
- claim escalation.

99. Narrative guard object

Every LLM generation call SHOULD receive:

```
NarrativeGuard

allowed_claims[]
required_claims[]

prohibited_claims[]

semantic_scope

confidence_tier

reference_scope

required_uncertainty[]

approved_terms[]
blocked_terms[]

maximum_claim_level

registry_version
```

100. Post-generation validation

Generated narrative SHALL pass a deterministic validator before display.

Validator checks at minimum:

```
unsupported_claim
identity_language
diagnostic_language
causal_language
deception_language
```

```
hidden_state_language
certainty_escalation
reference_misstatement
unresolved_overwrite
D3_inference
pattern_invention
interaction_invention
clinical_overreach
```

Failure:

```
NARRATIVE_REJECTED
```

and regeneration or fallback template.

101. Semantic entailment check

Every substantive sentence must be supported by one or more authorized semantic object IDs.

Recommended internal representation:

```
sentence_id
claim_text
supporting_object_ids[]
claim_level
confidence_tier
reference_scope
validation_status
```

A sentence with no supporting semantic object is not permitted.

102. Narrative fallback

If generative realization fails validation, SoulScope SHALL use deterministic approved templates.

Example:

```
fallback_template_id
```

rather than returning potentially unsafe free-form text.

103. Hallucination boundary

The Narrative Engine SHALL NOT invent:

- measurement values;
- Evidence;
- motivations;
- events;
- emotions;
- diagnoses;
- personal history;
- reasons for vocal change.

Unknown stays unknown in language exactly as it does in the scientific object.

104. Narrative Decision Record

Every user-facing result SHALL create:

```
NarrativeDecisionRecord

narrative_id
scan_id

source_semantic_result_id

pattern_source_id

dimension_source_ids[]
interaction_source_ids[]
evidence_source_ids[]

templates_used[]

claim_inventory[]
```

```
uncertainty_phrases[]

prohibited_claim_check

validator_results[]

model_id
model_version

prompt_template_version
narrative_registry_version

publication_status

created_at
```

105. Claim inventory

Every result SHALL maintain:

```
claim_id
claim_level
claim_text

source_ids[]

confidence_tier

scope

reference_scope

is_reflective_hypothesis

alternatives_present

validation_status
```

This permits exact traceability from prose back into reasoning.

106. Narrative provenance

For any displayed sentence, an authorized trace SHALL answer:

1. Which backend object supports this sentence?
 2. What claim tier is it?
 3. Which confidence tier applies?
 4. Which reference was used?
 5. Which uncertainty was preserved?
 6. Which template/model produced the language?
 7. Which prohibited-claim validator approved it?
-

107. Narrative acceptance tests

NAR-RAW-001

Given raw audio to Narrative Engine:

Expected:

ARCHITECTURE_VIOLATION

NAR-UNRES-001

Given:

Recovery = UNRESOLVED

Narrative:

"Your Recovery was balanced."

Expected:

REJECT

NAR-ID-001

Narrative:

"You are a restrained person."

Expected:

REJECT

NAR-DIAG-001

Narrative:

"Your voice suggests anxiety."

Expected:

REJECT

under current Canon.

NAR-CAUSE-001

Narrative:

"Higher Processing Demand caused your Stability to drop."

Expected:

REJECT

unless future validated causal authority explicitly permits it.

NAR-REF-001

First scan narrative:

"Higher than your usual pattern."

Expected:

REJECT

because no established personal baseline exists.

NAR-P1-001

Narrative:

"Compared with your neutral baseline..."

Expected:

REJECT

P1 is open reference, not neutral baseline.

NAR-P2-001

Narrative:

"During the negative emotion condition..."

Expected:

REJECT

NAR-P3-001

Narrative:

"Your voice recovered during the future-goals prompt."

Expected:

REJECT

NAR-MASK-001

Internal Interaction:

MASKS

Narrative:

"You were masking stress."

Expected:

REJECT

NAR-CONF-001

Backend:

Dimension high
confidence low

Narrative:

"This was clearly high."

Expected:

REJECT

NAR-NOPAT-001

Backend:

NO_PATTERN_PUBLISHED

Narrative creates a custom Pattern label.

Expected:

REJECT

NAR-TRACE-001

Every substantive sentence lacks support IDs.

Expected:

REJECT

108. User comprehension validation

Before narrative templates become production-active, studies SHALL evaluate whether users understand:

- that results apply to the scan;
- that acoustic evidence is not emotional certainty;
- that high confidence is not psychological certainty;
- that unresolved is legitimate;
- that population reference is not normality;
- that repeated Patterns are not personality;
- that Interaction verbs are not causal;
- that the user retains authority over meaning.

109. Comprehension failure

If a scientifically accurate phrase causes systematic misunderstanding, wording must change.

Scientific correctness is necessary.

It is not sufficient.

The communication must also be understood correctly.

110. Narrative localization

Localization SHALL preserve semantic force.

Translations must not strengthen:

may

into:

is

or:

suggested

into:

proved.

Localized templates require semantic review.

111. Accessibility

Narrative output SHALL remain understandable:

- without relying on color;
- without relying on the Resonance Signature image;
- with screen readers;
- under zoom/reflow;
- in plain text where required.

The semantic result must stand independently of visual rendering.

112. Relationship to Resonance Signature

The narrative may explain what the Signature represents:

“The Resonance Signature shows qualified acoustic structure across the timeline of your scan.”

It SHALL NOT say:

“This shape shows your emotional state.”

The narrative may connect both outputs to the same scan:

“The visual and written result come from the same recorded responses.”

But the narrative must preserve their distinct scientific roles.

113. Resonance Signature explanation

Approved conceptual language:

“The Resonance Signature is a time-resolved visual representation of qualified acoustic measurements from this scan. The written reflection comes from the separate Constellation reasoning system.”

This sentence may be used in educational/progressive disclosure material.

114. User edits and annotations

If users write notes about their result, these should be stored as:

USER_REFLECTION

separate from:

SOULSCOPE_INFERENCE.

User notes SHALL NOT silently enter future acoustic scoring unless a future explicitly adopted modality permits it.

115. Narrative versioning

A MAJOR Narrative Registry change is required when:

- allowed meaning changes;
- claim boundaries change;
- a new claim tier is introduced;
- a user-language key changes scientific interpretation;
- Pattern/State/Interaction wording materially changes semantic force.

Copyediting without semantic change may be MINOR/PATCH according to active governance.

116. Historical narrative

Historical narrative results SHALL retain:

```
narrative_registry_version
template_version
model_version
semantic_result_version
```

A later wording change does not silently rewrite what was previously shown.

117. Narrative regeneration

If a user asks to view an old result under newer wording, the system may create:

```
ORIGINAL_NARRATIVE
REGENERATED_NARRATIVE
```

provided both remain tied to the same original Semantic Result unless scientific reprocessing is separately authorized.

118. Hard invariants

NAR-INV-01

Narrative does not reason.

NAR-INV-02

Narrative does not score.

NAR-INV-03

Narrative cannot strengthen upstream certainty.

NAR-INV-04

Narrative cannot resolve missing data.

NAR-INV-05

Narrative cannot infer D3 constructs.

NAR-INV-06

Narrative cannot create Pattern, State, or Interaction labels.

NAR-INV-07

Narrative must preserve reference identity.

NAR-INV-08

Narrative must preserve scan scope.

NAR-INV-09

Narrative cannot turn correlation/coexistence into causation.

NAR-INV-10

Narrative cannot turn acoustic evidence into emotion certainty.

NAR-INV-11

Narrative cannot turn repeated observation into personality.

NAR-INV-12

Narrative may offer reflective hypotheses only conditionally.

NAR-INV-13

Every substantive claim must trace to an authorized semantic object.

NAR-INV-14

A deterministic validator must check generated prose before publication.

NAR-INV-15

When language cannot safely express the result, simpler language or no claim is preferable.

119. What v0.1 freezes

The following are canonical:

- narrative is downstream translation only;
 - completed Semantic Result is the sole semantic source;
 - scope-first language;
 - observation before interpretation;
 - separate confidence and magnitude;
 - explicit unresolved language;
 - first-scan versus personal-baseline distinction;
 - P1 is open reference, not neutral baseline;
 - P2/P3 are contexts, not emotion ground truth;
 - Dimension-specific semantic boundaries;
 - safer Interaction translations;
 - Pattern describes scan, not person;
 - Tier C reflection remains optional and conditional;
 - claim-level architecture;
 - claim provenance;
 - LLM guard object;
 - deterministic post-generation validation;
 - fallback templates;
 - historical narrative versioning;
 - Resonance Signature and narrative remain scientifically distinct.
-

120. What v0.1 intentionally does not freeze


The following still require user research:


- exact final user-facing phrasing;
- preferred reading level;
- Pattern headline wording;
- whether some internal Interaction verbs should ever appear directly;
- confidence-band wording thresholds;
- quantity of technical detail users prefer;
- reflection-prompt style;
- localization language;
- final result length;
- progressive-disclosure depth;
- comprehension thresholds.

Those should be determined through comprehension testing rather than aesthetic preference alone.

121. Complete semantic pipeline

The entire SoulScope semantic system is now:





with:

Evidence Ledger + Decision Ledger + Narrative Provenance

preserved throughout.

122. Final Narrative law

SoulScope should never make the result sound more certain, more psychological, more causal, or more personal than the evidence beneath it.

The user-facing language should become simpler as it moves upward.

The science underneath should remain fully reconstructable.

Measure precisely → reason conservatively → communicate clearly

And at the final boundary:

SoulScope explains what it noticed. The individual remains the authority on what it means.